

AI Factories* (SW)

Forkable Sandboxes as a Runtime for Automated Software Development

Yossi Eliaz

Principal Engineer, Incredibuild / islo.dev

Associate Professor, Holon Institute of Technology (HIT)


Computer Laboratory Systems Research Group Seminar
University of Cambridge

15 October 2026 - FW11 + Microsoft Teams
30-minute version

Burn the runner. Keep the proof. Fork the machine, not the trust.

* “Factory” means an atelier-like search workshop, not an assembly line.

100 agents. One laptop. What could possibly go wrong?



Parallelism without isolation is correlated chaos.

Why I ended up here

Computational science

An experiment is not reproducible unless the state, inputs, and provenance are reproducible.

Production ML + builds

Reuse everything you safely can: caches, artifacts, warm state. Rebuilding the world is expensive.

Coding agents

Then the developer became a program that edits files, runs shells, calls networks, retries, and forks hypotheses.

The collision

The same **state** I wanted to reuse for speed became the state I had to isolate for trust.

That is the road from reproducible science and build systems to **forkable sandboxes**.

Search can be plural. Authority must be singular.

cheap and plural

scarce and singular



100 candidate patches

1 merge to main

10,000 rollouts

1 policy checkpoint

millions of sim trials

1 robot / wet-lab hour

SEARCH: fail cheaply

AUTHORITY: commit carefully

The runtime should make the left side cheap *without letting it become the right side.*

There are two ways to make futures cheap

Learn a world

Dreamer / Contrastive World Models

$$o_t \rightarrow z_t \rightarrow \hat{s}_{t+1}, \hat{s}_{t+2}, \dots$$

- ▶ compact, learned latent state
- ▶ imagine many futures cheaply
- ▶ useful when interaction is slow, unsafe, or scarce
- ▶ pay approximation / learned-dynamics error

CWM's twist: stop spending capacity reconstructing irrelevant pixels; keep information predictive of the future.

Fork the world


executable software environment

$$S_0 \rightarrow \{S_0^1, S_0^2, \dots, S_0^N\}$$

- ▶ explicit, executable state
- ▶ branch real repo / process futures
- ▶ useful when the world is software and reset is cheap
- ▶ pay compute / storage, not learned transition-model error

Exact only inside the state and capability boundary we actually captured.

For code, the best world model is often the world



These are **executable counterfactuals**: branch from the same past, intervene differently, observe the actual software transition.

No learned transition model is required for the part of the world we can execute faithfully.

Related systems: Firecracker (NSDI'20), DeltaBox (arXiv:2605.22781), Shepherd (arXiv:2605.10913).

A fork should copy state, not authority

Copy into the child

- ▶ filesystem + process / memory state
- ▶ pinned source, dependencies, toolchain
- ▶ enough state to continue the experiment

Share carefully

- ▶ content-addressed build artifacts
- ▶ immutable heat / caches inside one trust domain
- ▶ CoW pages whose sharing policy is explicit

Remint at the boundary

- ▶ API / cloud credentials
- ▶ network capabilities and identity
- ▶ randomness / epoch identity when required

Keep outside the child

- ▶ hidden verifier + objective digest
- ▶ promotion key / merge authority
- ▶ durable evidence ledger

Coding RL turns reset into a training primitive

For coding RL, the environment is part of the training example.

$$\text{episode} = (S_0, \text{task}, \text{tools}, \text{verifier}, \text{reset})$$

Cold path

clone repo → install deps → build →
run episode → throw away

Reconstruct the same starting world again.

Fork path

prepare once → snapshot S_0 → fork
episode → burn child → fork again

Reset becomes a runtime operation.

Systems thesis

Reset is part of the reward contract: episodes are comparable only if they begin from an authenticated S_0 .

SWE-smith: [arXiv:2504.21798](https://arxiv.org/abs/2504.21798). Meta Code World Model (2025). RL-environment products: Scale, Surge, Mercor, Prime Intellect.

The student cannot grade their own exam

Inside the child

code, shell, files, compiler, tests it is
allowed to see

untrusted searcher

Outside the child

hidden verifier, objective digest,
reward policy, promotion key

trusted authority

RUN $\neq$ EVAL

If the agent can rewrite the ruler, a green score stops meaning what you think it means.

Reward-hacking example: From Rebound to Remedy, arXiv:2604.01476.

Four clones do not make four witnesses

$$S_0 \rightarrow \{w_1, w_2, w_3, w_4\}$$

They may share:

- ▶ the same prompt and model
- ▶ the same cached artifact
- ▶ the same hidden bad fixture
- ▶ the same ancestor state

So a worker returns more than a score
result + precision + evidence IDs +
lineage

Cold liar

If a worker can claim absurd confidence,
it can dominate a reducer even when it is
wrong.

**Fork makes search cheap. It does not manufacture
independent evidence.**

Evidence-Aware / Boltzmann MapReduce: Eliaz, arXiv:2607.09689.

Three factories. Same loop.

snapshot -> fork -> search -> verify -> reduce -> promote ->
burn


Factory	Search	Oracle / reduce	Scarce tip
Software	candidate patches	tests + evidence lineage	merge to main
RL training	rollout trajectories	reward / verifier	policy checkpoint
HIL / science	sim / hypotheses	measured evidence	robot / wet / paper

The isomorphism

Same runtime contract; different objectives.

Jointly test compositions; promotion gets costlier toward production and science.

The capability contract is only three verbs



Durable intent. Disposable execution. Singular authority. Deterministic convergence.

Search may be stochastic. **Authority may not be ambient.**

What I want from this room

Four systems questions I do not think we have answered cleanly yet:

1. **Executable state:** what must a first-class machine value include – FS, memory, process tree, GPU, sockets, clock, external state?
2. **Model vs world:** where is the crossover where learned imagination becomes cheaper than forked executable interaction?
3. **Correlation:** how should reducers price shared ancestry instead of pretending siblings are i.i.d.?
4. **Authority:** how do children inherit enough state to act without inheriting the power to grade or promote themselves?

If you remember one thing

Burn the runner. Keep the proof. Fork the machine, not the trust.

Selected sources / Q&A guardrails

World models	Hafner et al., <i>Training Agents Inside of Scalable World Models</i> (Dreamer 4), arXiv:2509.24527; Li, <i>Contrastive World Models</i> , arXiv:2609.22175.
Fork systems	Firecracker, NSDI 2020; DeltaBox, arXiv:2605.22781; Shepherd, arXiv:2605.10913.
Evidence-aware reduce	Eliaz, <i>Evidence-Aware MapReduce / Boltzmann MapReduce</i> , arXiv:2607.09689.
Reward integrity	<i>From Rebound to Remedy</i> , arXiv:2604.01476.
Training environments	SWE-smith, arXiv:2504.21798; Meta Code World Model (2025); Prime Intellect; Scale / Surge / Mercor environment products.

Claim fence

A fork removes learned-dynamics error only inside the executable boundary it actually captures. External services, nondeterminism, hardware and incomplete tests can still make the world partial. Fork also does not imply independent evidence.