

# Forkable Sandboxes

## The Runtime Layer for AI Software Factories

Forks multiply executions. Who counts the evidence?

**Yossi Eliaz**

Principal Engineer, Incredibuild · Associate Professor, Holon Institute of Technology (HIT)

Disclosure: I build an agent sandbox platform. I name components; I do not rank vendors.



# Hedging: ask twice, keep the first answer

If a server is slow, ask a second server too, and use whichever answers first.

**1,800 ms  $\rightarrow$  74 ms**

Google's slowest 0.1% of requests. Cost: 2% more requests.

“the source of latency is often not inherent in the particular request”: the server is slow, not the request.

# Hedging: ask twice, keep the first answer

If a server is slow, ask a second server too, and use whichever answers first.

**1,800 ms  $\rightarrow$  74 ms**

Google's slowest 0.1% of requests. Cost: 2% more requests.

“the source of latency is often not inherent in the particular request”: the server is slow, not the request.

**Agents hedge for correctness. A flaw in the task repeats in every copy.**

Dean & Barroso, The Tail at Scale, CACM 2013: backup after 10 ms; 99.9th-percentile latency reading 1,000 BigTable keys. It works only when the slowdown does not hit several servers at once. Shared mistakes across different AI models: [Kim et al., ICML 2025](#).

# Forks find answers. Checks decide which.

# 16% → 56%

SWE-bench Lite, one sample vs. 250  
(Brown et al. 2024)

- ✓ with an automatic check, more samples solve more problems
- ✗ without one, voting and reward models plateau
- ✗ a check that says yes to wrong answers caps the gain; when a wrong answer costs more than none, the best number of tries is often under ten (Stroebl et al. 2024)

**Hedging for correctness is only as good as the check you hedge on.**

Brown et al., Large Language Monkeys: Scaling Inference Compute with Repeated Sampling, arXiv:2407.21787;  
Stroebl, Kapoor & Narayanan, The Limits of Inference Scaling Through Resampling (v1: Inference Scaling fLaws), arXiv:2411.17501.

# Numbers every agent runtime should know


| Operation                                            | Number           | Source                           |
|------------------------------------------------------|------------------|----------------------------------|
| checkpoint an agent sandbox                          | 10.8 ms          | DeltaBox, 2026                   |
| resume a paused microVM                              | 49 ms, best case | Kimi K3, 2026                    |
| boot a microVM to application code                   | <125 ms          | Firecracker, NSDI'20             |
| fork an agent sandbox                                | 134–143 ms       | Shepherd, 2026                   |
| fork a Xen VM across hosts                           | 600–800 ms       | SnowFlock, EuroSys'09            |
| restore, run, capture through a public API           | 6.87 s p50       | Eliaz, arXiv:2607.09689, Table 2 |
| one software-factory work order                      | 43 min, \$10.33  | my run, 27 Sep 2026              |
| a sandbox's life spent waiting on the model          | up to 98%        | Kimi K3, 2026                    |
| sandboxes across one model's training and evaluation | 51.2 million     | Kimi K3, 2026                    |

how correlated two forks' verdicts are ? no surveyed system reports it

## Every row has a source, except the one that matters most.

DeltaBox; Kimi K3 report §5.3.2; Agache et al. NSDI'20; Shepherd; Lagar-Cavilla et al. EuroSys'09; arXiv:2607.09689; our Airflow Factory runs, Sep 2026 (time, cost). Survey: slide 22. After J. Dean, "Numbers everyone should know" (LADIS 2009).

# Where 43 minutes went



Make sandbox setup a hundred times faster, and this run gets 0.8% faster.


**ANALOGY** Airflow Factory is self-driving software: an issue goes in, a reviewed patch comes out.

**The minutes are in the agents and the checks. So is the risk.**

Our Airflow Factory runs, Sep 2026: stage timings (one cycle 2,557 s; sandbox setup 20.3 s).

# Blocker. One more repair. Approve.

Work order: fix the sandbox allowlist. A four-thread race test failed; the loop sent a repair agent.




One model, four steps. All 1,972 tests passed after each repair. The new code had no test of its own.

**Both repairs said they had not run the tests. The gate approved.**

# Blocker. One more repair. Approve.

Work order: fix the sandbox allowlist. A four-thread race test failed; the loop sent a repair agent.



One model, four steps. All 1,972 tests passed after each repair. The new code had no test of its own.

**Both repairs said they had not run the tests. The gate approved.**

Suppose nine forks of repair 1 all pass. How many witnesses is that?

Our Airflow Factory runs, Sep 2026: the agent records of both repairs and both reviews.

# Why a biological physicist forks sandboxes


The same question, met six times: how many independent witnesses do I really have?



**Six fields, one bug: evidence that looks independent and is not.**

Physics and biology are vocabulary here, not evidence about forks. Genomics and Bitcoin are team results with my co-authors. Physics: PRE 2020, PNAS 2020, JPCB 2021; genomics: Hitz et al. 2023, iScience 2026; Bitcoin: arXiv:2206.02871; self-driving: my perception work at Mobileye, 2021–2024; factories: our Airflow Factory runs, arXiv:2607.09689; RL: Kimi K3 report §5.3.2.

# One factory, seven places the runtime bit



**MEASURED** my run records    **BUILT** built and tested, or synthetic    **PROPOSED** design, not yet run  
**PUBLISHED** other people's work    **ANALOGY** borrowed vocabulary, not evidence

**Scope, said once:** two Airflow Factory runs, in Docker Sandboxes through Airflow's sandbox toolset, models served through Databricks; two observations from my own sandbox platform; my harness answered the gates; leases off; nothing forked.

# Intelligence: choose between, then act

**intelligence**, from Latin *intellegere*: *inter* (between) + *legere* (to choose, to read).

Intelligence is choosing between.


Since 29 Sep 2026, US federal agencies must call AI 'Super Intelligence' (SI). I will just say intelligence.

**ANALOGY** My ladder, as I frame it:

**self-driving cars**: sense, plan, act (Mobileye)

**self-driving computers**: software that writes and tests itself (Airflow Factory)

**self-driving intelligence**: an agent in a loop




**act** happens in a sandbox: **wall**. **choose** needs alternatives: **fork**. **check** decides what counts: **count**.

**Intelligence acts on what it reads. The sandbox is where it acts.**

Oxford English Dictionary / Online Etymology Dictionary, 'intelligence'. Executive order, 29 Sep 2026 (CNBC, The Register, 29–30 Sep 2026). Mobileye: my perception work, 2021–2024.

# A factory needs a wall, a fork and a count



## I claim

**BUILT** a receipt contract: reused evidence stops the merge;  
every result carries its lineage, a record of where it came from

**MEASURED** two real factory runs

**PROPOSED** H1: snapshot restore beats a warm cache by  $>10$  s


**PROPOSED** H2: siblings from one snapshot fail together more  
than strangers

 **Forks multiply executions, not evidence.**  
**The runtime sees what forks share.**

## I do not claim

- ✗ a new estimator
- ✗ a controlled speedup
- ✗ a ranking of platforms
- ✗ a factory that forks today
- ✗ physics or biology as evidence about forks

# Allowlists name hosts; the web redirects




- the list lived in five places: two scripts, a checker, two docs
- the fix added both hosts
- the second one admits every release asset on GitHub

**Fetch by digest through a caching proxy, not by name.**

Observed on my sandbox platform, 27 Sep 2026. Work order: our first Airflow Factory run, Sep 2026.

# Keys stay outside; a fork copies the inside



**PROPOSED** Next: the gateway keys each request on host-assigned identity (per-clone tap or vsock CID), never on a string in guest memory. Restore bumps a generation and voids the parent's lease.

## Copy the machine, not the keys.

Airflow Factory design notes on sandbox isolation and credential authority; [Docker Sandboxes security](#); Capsicum (USENIX Sec'10); macaroons (NDSS'14); Gray & Cheriton, Leases (SOSP'89).

# Filesystem state: what a snapshot shares

| State                | Policy                       | Known mechanisms                                                                                       |
|----------------------|------------------------------|--------------------------------------------------------------------------------------------------------|
| files, build outputs | pin by content digest        | OBuilder: CoW snapshot per build step (ZFS, btrfs)                                                     |
| workspace disk       | copy-on-write layers         | overlay layers (day10); DeltaBox's layered CoW filesystem; chained overlaybd snapshots (DeepSeek DSec) |
| agent history        | keep outside the VM          |                                                                                                        |
| randomness           | reseed, or couple on purpose | later: slide 38                                                                                        |
| external services    | mediate                      | later: slide 27                                                                                        |

**A snapshot captures a boundary, not the universe.**

OBuilder; day10; DeltaBox; DeepSeek-V4 §5.2.5.

# Locking the tests was not enough

## In my run

- ✓ tests/ write-protected, and the agent was told
- ✓ the test hook never fired: 15 edits, all allowed
- ✗ planned steps are scope-checked; repairs are not
- ✗ a repair edited code outside the plan
- ✗ an out-of-plan edit is “major”, not a blocker
- ✗ writer and both reviewers: one model

Two fixes: scope-check repairs; make an out-of-plan edit a blocker.

## A locked directory is not a scope.

Our Airflow Factory runs, Sep 2026: repair 1's record (an edit outside the plan); the factory's repair stage records repairs without a scope check. Barkai, . . . , Eliaz, . . . , OffRisk, *Bioinformatics Advances* 3:vbad138 (2023).

### ANALOGY · BIOLOGY

Cas9 can cut where the guide did not point: an **off-target** site.

OffRisk, a Docker image my co-authors and I built, annotates CRISPR off-target sites and labels their risk (2023).

A scope check is that label for code.

# Reward hacking is the RL name for it

| Source          | What it found                                                                                                        |
|-----------------|----------------------------------------------------------------------------------------------------------------------|
| METR (2025)     | o3 reward-hacked in <b>30.4%</b> of RE-Bench runs vs. <b>0.7%</b> of HCAST; METR's leading guess: the visible scorer |
| ImpossibleBench | read-only tests stop test edits, not special-casing                                                                  |
| Kimi K3 §4.2.6  | isolate agents from verifiers; pair public diagnostic verifiers with hidden ones                                     |

**Kimi K3, kernel-optimisation tasks (§4.2.4).** Its hacking-detection system penalises “CUDA graph replay, input caching, and precision reduction”.

In a build farm, caching inputs is the job. Under a speed grader, it is the hack. The grader decides which.

**Run  $\neq$  evaluate  $\neq$  promote: three permissions, not one.**

METR (June 2025); ImpossibleBench; Kimi K3 technical report §4.2.4, §4.2.6.

# Teardown must not destroy the experiment



The stray file was my own harness's review-diff archive, mirrored into the workspace.

Our second work order: 26 min, \$3.43, nothing delivered.

## The gate was right. Recovery was wrong.

The artifact leaves the cell before the cell dies. In a long human wait: snapshot, destroy, restore.

Our second Airflow Factory run, Sep 2026: run records (26 min wall clock, 853.8 s stage time, \$3.43); redacted copy on request; fix not yet merged. Related: Planarian statepoints (arXiv:2609.35366); Crab turn-level recovery (arXiv:2604.28138); DeepSeek-V4 DSec keeps the sandbox when its training task is preempted and replays cached results on resumption (§5.2.5).

# A green check is not proof the check ran

## Sandbox eval CI job

`success` in 3 s; every real step skipped, because the key was unset

## Approval record

```
"mode": "human", actor  
admin; my harness answered
```

## My sandbox CLI

status lines mixed into stdout; an earlier attempt died parsing them

All three are my own receipts, and each was wrong.

## Every check needs a receipt the sandbox cannot write.

**RL learned the same thing.** Kimi K3 grounds reward in the verifier's evaluation of the final environment state "rather than the agent's self-reported completion" (§4.2.6).

Our CI, 28 Sep 2026; our Airflow Factory runs, Sep 2026: the approval record and both repair records. Provenance: PASS (USENIX ATC'06), CamFlow (SoCC'17), in-toto (USENIX Security'19). [Kimi K3](#) §4.2.6.

# What the run pins, and what it cannot

## Pinned in the run record

- ✓ input commit (8ff967d) and `uv.lock`
- ✓ policy and inputs digests
- ✓ sandbox adapter digest
- ✓ every agent step's output, as JSON

## Not pinned

- ✗ the hosted model: no sampling seed
- ✗ external hosts and their redirects
- ✗ the clock and scheduling
- ✗ the adapter source: a local shim, not in git


**PROPOSED** Replay, not rerun: log every model response, external fetch and clock read; re-execute the rest against the log.

Genomics set this bar (slide 7): each ENCODE data file is captured with the software versions, parameters and reference genome that produced it.

## You can audit this run. You cannot rerun it.

Our Airflow Factory runs, Sep 2026: stage timings and work-graph execution record. [Hitz et al., The ENCODE Uniform Analysis Pipelines, bioRxiv 2023](#). Record and replay: `rr` (USENIX ATC'17); DeepSeek-V4 DSec logs every sandbox command and result for replay ([arXiv:2606.19348 §5.2.5](#)).

# Parallel-safe steps, and nothing to fork them



reason: `serial_fallback_missing_fork`

The steps touch different files, so two worktrees would do here. A fork matters when the reached state is expensive.

**Three minutes of forty-three. We fork to buy alternatives.**

Our Airflow Factory runs, Sep 2026: work-graph execution record (node timings measured; the fork lane is computed, fork cost excluded). Remote fork at scale: MITOSIS (OSDI'23).

# What our runs found

---

- |                                              |                                                        |
|----------------------------------------------|--------------------------------------------------------|
| <b>1 The sandbox is not the bottleneck</b>   | setup was under 1% of a 43-minute work order           |
| <b>2 A locked test folder is not a scope</b> | a repair rewrote code outside the plan                 |
| <b>3 Green checks can be empty</b>           | CI skipped every step; a review approved untested code |
| <b>4 One model did all four steps</b>        | wrote, reviewed, fixed, approved: one look, not two    |
| <b>5 Teardown can destroy the result</b>     | an approved patch died with its sandbox                |
| <b>6 Nothing forked</b>                      | the factory asked for a fork; none was available       |
- 

**Seven surfaces, one rule: whatever you must trust lives outside the cell.**

Our Airflow Factory runs, Sep 2026: two work orders in Docker Sandboxes, models served through Databricks.  
Records: [github.com/zozo123/ariflow-swfactory](https://github.com/zozo123/ariflow-swfactory).

# The shape of an agent sandbox, drawn in 2009

## (a) Sandboxing

```
state = trusted_code()
ID = VM_fork(1)
if ID.isChild():
    untrusted_code(state)
    VM_exit()
else:
    VM_wait(ID)
```

## (b) Parallel Computation

```
ID = VM_fork(N)
if ID.isChild():
    parallel_work(data[ID])
    VM_exit()
else:
    VM_wait(ID)
```

In (b), `data[ID]` makes each child different. Fork an agent's repair nine times: nine children, one slice.  
Next: from Xen (2003) to training gyms (2026).

**The mechanism is old. The caller is new: a program that searches.**

Lagar-Cavilla et al., SnowFlock: rapid virtual machine cloning for cloud computing, EuroSys'09, Fig. 1(a) and 1(b); built on Xen 3.0.3 (Barham et al., SOSP'03).

# Fork is old, and much of it was built on Xen

| Year | System |
|------|--------|
|------|--------|

|      |                 |
|------|-----------------|
| 2003 | Xen (Cambridge) |
|------|-----------------|

|      |                            |
|------|----------------------------|
| 2005 | Live migration (Cambridge) |
|------|----------------------------|

|      |          |
|------|----------|
| 2005 | Potemkin |
|------|----------|

|      |           |
|------|-----------|
| 2009 | SnowFlock |
|------|-----------|

|      |           |
|------|-----------|
| 2020 | Catalyzer |
|------|-----------|

|      |                  |
|------|------------------|
| 2023 | Nephele; MITOSIS |
|------|------------------|

| Year | System |
|------|--------|
|------|--------|

|      |          |
|------|----------|
| 2026 | DeltaBox |
|------|----------|

|          |
|----------|
| Shepherd |
|----------|

|      |
|------|
| Crab |
|------|

|           |
|-----------|
| Planarian |
|-----------|

|                             |
|-----------------------------|
| Kimi K3 AgentENV (training) |
|-----------------------------|

**Twenty-three years of making fork cheap. No surveyed system counts what siblings share.**

SOSP'03; NSDI'05; SOSP'05; EuroSys'09; ASPLOS'20; EuroSys'23; OSDI'23; [DeltaBox](#); [Shepherd](#); [Crab](#); [Planarian](#); [Kimi K3](#) §5.3.2.

# Training runs are sandbox factories

## 51,219,741

sandboxes, across 1,505,678 images

Kimi K3, training + evaluation

- **pause** while the model thinks: up to 98% of a sandbox's life, using no memory or CPU
- **fork** from the exact state, original still running: “for reward judging without side effects”
- **snapshot** at regular intervals, for error recovery
- **partial rollout**: optimise once a fraction  $\lambda$  of trajectories completes, “to mitigate the long-tail latency”

DeepSeek-V4's DSec: hundreds of thousands of concurrent sandboxes in one cluster.

Partial rollouts are the same move as Dean & Barroso's “good enough”: stop waiting for the stragglers.


**A sandbox is an environment; a fork is a reset; a verifier is a reward.**

Kimi K3 §4.1.2, §5.3.2 (AgentENV); DeepSeek-AI, DeepSeek-V4, [arXiv:2606.19348](#) §5.2.5; Dean & Barroso, The tail at scale, CACM 2013.

# Reset changes the learning problem

Where each attempt starts decides what the learner sees.

$$s_0 \sim \mu \quad \text{versus} \quad s_0 \sim \mu_{\text{checkpoint}}$$



Go-Explore: first return, then explore (Ecoffet et al., Nature 2021)

- **PROPOSED** record each checkpoint's origin and the rule that selected it
- **PROPOSED** evaluate the final policy on the intended task distribution

Same for a factory that forks after the plan: each child's success is conditional on that plan (slide 47).

## A restore point is a choice of training data.

Ecoffet et al., First return, then explore, Nature (2021).

# Fork copies memory, secrets and identity

| The child inherits   | Hazard                        | What is known                                                         |
|----------------------|-------------------------------|-----------------------------------------------------------------------|
| random-number state  | siblings share random streams | VMGenID reseeds the kernel only; userspace has “no generic solution”  |
| tokens in memory     | $N$ live copies               | keep secrets outside; wipe on suspend (Brooker et al. 2021)           |
| guest clock          | resumes at snapshot time      | resync on restore                                                     |
| IP and MAC           | identical identities          | namespace + NAT per clone; Potemkin spent 142.8 of 521 ms on IP setup |
| open TCP connections | cannot be cloned              | CRIU hands one to one restore; gVisor returns ECONNRESET              |

## Eight forks of a VM holding a token are eight live tokens.

Firecracker snapshotting docs; [Brooker et al.](#), [arXiv:2102.12892](#); Vtable et al. SOSP’05; CRIU and gVisor checkpoint/restore docs.

# What each kind of fork copies

$M_{\text{total}} \approx M_{\text{shared}} + \sum_i M_{\text{private},i}$ : children share until they write.

Kimi K3, real workloads: copy-on-write memory plus page-cache optimisations, up to  $6.5\times$  overcommit.


| Mechanism                   | What survives                                                                     | Where the cost lands                      |
|-----------------------------|-----------------------------------------------------------------------------------|-------------------------------------------|
| worktree, overlay layer     | files only                                                                        | bytes each child writes                   |
| container checkpoint (CRIU) | processes and memory; a TCP connection needs its original IP, so one copy at most | dump size; what the kernel can checkpoint |
| microVM snapshot            | guest memory and vCPU/device state; disks handled separately                      | pages each child touches after restore    |
| record and replay (rr)      | a log, not a state                                                                | replay time                               |

A restore that returns OK is not a faithful copy: compare restored and direct runs on declared outputs before counting a child.

## Cheap at fork; you pay at divergence.

Firecracker snapshot-support docs; CRIU TCP\_REPAIR docs; REAP (ASPLOS'21); rr (ATC'17); Kimi K3 §5.3.2.

# Hold effects until an authority releases them



- Speculator, external synchrony, Remus: run ahead, hold output until it is safe
- Zheng et al. (2026, stated in Lean): an execution edit cannot undo a tool request already sent
- DSec replays cached results on resume, so a command unsafe to repeat never runs twice
- **BUILT** here: children hold no publishing credential; the orchestrator delivers after the gate
- exception: the model API is an unheld channel, metered and logged

## Output commit, made structural.

Speculator SOSP'05; external synchrony OSDI'06; Remus NSDI'08; [Zheng et al., arXiv:2608.22928](#); DeepSeek-V4 §5.2.5.

# When fork pays, and where it loses

$$(N - 1) P > H + N (R + D)$$

saved:  $N - 1$  preparations  $P$       spent: one capture  $H$ , plus a restore  $R$  and a divergence  $D$  per child  
Work  $W$  (model calls included) and evaluation  $V$  cancel: each child still rebuilds its patch.

## Where fork loses

- ✗  $N = 1$
- ✗ the state is only files and the build cache is good
- ✗ a minimal image boots faster (Jitsu; LightVM)


**If your build cache is good, fork loses. Fork pays for state you cannot rebuild.**

Eliaz, [arXiv:2607.09689](https://arxiv.org/abs/2607.09689), Table 2; Madhavapeddy et al. NSDI'15 (Jitsu); Manco et al. SOSP'17 (LightVM, NEC Labs).


# Isolation has two axes: kernel and authority

## Kernel surface

less shared



## Authority



**PROPOSED** Warm state, shared on purpose: build cache within one trust domain (risk: timing and content leaks) · dependency image by digest (risk: a poisoned base) · scratch never (risk: cross-branch contamination)

**Pick the kernel surface for the workload. Hand over authority explicitly.**

**Kimi K3** §5.3.2; Firecracker production host setup guide; Watson et al., Capsicum (USENIX Sec'10) and CHERI (IEEE S&P'15).

# The interface the runtime owes the factory


| Call                                                                                                     | Status                                                                               |
|----------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------|
| <code>checkpoint(cell) → snap</code>                                                                     | <b>MEASURED</b> my platform's named snapshots; memory unverified                     |
| <code>fork(snap, lease, seed=copy reseed, egress)</code>                                                 | <b>PROPOSED</b>                                                                      |
| <code>evaluate(artifact, verifier) → receipt</code>                                                      | <b>PROPOSED</b> runs outside the child                                               |
| <code>select(candidates, fresh test) → one</code>                                                        | <b>PROPOSED</b> choosing, not pooling                                                |
| <code>reduce(receipts) → (<math>\hat{\theta}</math>, <math>N_{\text{eff}}</math> bound)   abstain</code> | <b>BUILT</b> merge + evidence check; <b>PROPOSED</b> $N_{\text{eff}}$ bound, abstain |
| <code>promote(epoch, receipt)</code>                                                                     | <b>BUILT</b> sha256-bound gate; <b>PROPOSED</b> epoch fence                          |

The epoch fence stops a stale run from writing: a Chubby sequencer or a fencing token.

## Choosing is not pooling: select and reduce are different calls.


Airflow Factory fork contract; my platform's named snapshots: [arXiv:2607.09689](https://arxiv.org/abs/2607.09689) §4; Burrows OSDI'06; Kleppmann (2016). Architecture: next slide.

# Machines fork; authority stays outside




Airflow Factory design docs: fork lifecycle and run recovery. Solid: built and ran. Dashed gold: not exercised (lease broker built; evaluator proposed). Dashed red: missing.

# Machines fork; authority stays outside



Airflow Factory design docs: fork lifecycle and run recovery. Solid: built and ran. Dashed gold: not exercised (lease broker built; evaluator proposed). Dashed red: missing.

# Machines fork; authority stays outside



Fork contract, written before the fork (all eight conditions in the end matter)

5. never count siblings that reuse the same evidence as independent agreement

8. tear down every child, on any exit

our second run

I wrote the fork contract before I had the fork.

Airflow Factory design docs: fork lifecycle and run recovery. Solid: built and ran. Dashed gold: not exercised (lease broker built; evaluator proposed). Dashed red: missing.

# Promotion is small enough to model-check

A TLA+ design model of one cell, from epoch replacement to publication.

- evidence refers only to the frozen candidate
- approval requires evidence for that same candidate
- publication requires frozen candidate = evidence = approval under the current Airflow-bound epoch
- an external commit belongs to the current epoch
- a search-origin authority request cannot become an authority source
- an `in_doubt` effect has no blind retry and leaves that state only by observation

**Design model: no TLC model-checker run and no trace check recorded yet. The epoch fence is proposed.**


Dean & Barroso: tail-tolerant techniques apply most to operations that do not perform critical mutations; consistent updates use quorum algorithms such as Paxos.

**Model-check the gate; count the children.**

Airflow Factory TLA+ authority model; Dean & Barroso, The tail at scale, CACM 2013 (“Mutations”).

# The second review was not a second witness

Nine strangers agree: probably right. Nine with one wrong map: one opinion.



Repair 2 answered review 1's verdict, and the same model graded it: a reused holdout (Dwork et al. 2015).

**Choosing** among 9 different patches needs a fresh test;  $\text{best-of-9} \leq 1 - P(\text{all fail})$ .

**Confirming** one patch 9 times is capped by what the runs share.

**ANALOGY · BIOLOGY** In our 2021 CRISPR-screen paper, each viability point was at least four technical replicates; each plot, at least two independent experiments.

## Nine forks of one parent are technical replicates.

Our Airflow Factory runs, Sep 2026: the records of review 1, repair 2 and review 2. Dwork et al., The reusable holdout, Science 2015. [Chen, arXiv:2606.27288 \(2026 preprint\)](#). Gilad, Eliaz et al., Commun. Biol. 4:399 (2021), Methods.

# A reused grader becomes a training set

Even a protected evaluator can teach the search to overfit its feedback.



## STATISTICS

A reused holdout overfits (Dwork et al. 2015); bound the bias of adaptive analysis (Russo & Zou 2016).

## RL TRAINING

Kimi K3: public verifiers give diagnostic feedback; hidden ones score held-out scenarios. A verbosity cap curbs reward hacking by length.

**PROPOSED** For agents: grader outside the child, every query logged, freeze before the final test.

## Feedback you optimise against stops being evidence.

Russo & Zou, AISTATS 2016; Dwork et al., Science 2015; Kimi K3 report §4.2.6 (verifiers), §4.1.2 (verbosity). Visible scorer: slide 15.

# One formula, three fields

$$P(\text{at least one of } n) = 1 - (1 - p)^n$$

## SYSTEMS

Dean & Barroso: each server is slow on 1 request in 100 (99th percentile 1 s); a request fans out to 100 servers.

# 63%

of requests are slow.

## BIOLOGY

Dorfman 1943: a pool tests clean only if every sample is. Poolkeh, two-step pools for  $\sim 9$  M people, would need

# 295,951

tests: 30-fold fewer (a model, not a deployment).

## THIS TALK

Suppose the race breaks 1 run in 10. Green runs needed before a miss is under 5%:

# 29

because  $0.9^{29} = 0.047$ .

## All three assume independent draws.

Dean & Barroso, The tail at scale, CACM 2013. Dorfman, Ann. Math. Stat. 14:436 (1943). Eliaz, Danovich & Gasic, Poolkeh, medRxiv 2020 (Israel,  $s_1 = 92$ ,  $s_2 = 10$ ). Race: our first run, review 2; the 1-in-10 rate is illustrative.

# More forks are not more witnesses

100 forks that share a little: 9 witnesses. A lot: 2. Identical: 1.

$$\text{Var}(\bar{Y}_N) = \sigma^2 \left( \rho + \frac{1 - \rho}{N} \right)$$


$$N_{\text{eff}} = \frac{N}{1 + (N - 1)\rho}$$

$\rho$ : how much the copies share.

A little = 0.1, a lot = 0.5, identical = 1.

4 re-runs  $\times$  100 tests = 400 executions, 100 evidence IDs. Pool within a candidate; compare across.

Different models, both wrong: same answer 60% (chance 33%). 9 LLM judges: 2.2 independent votes.



Equal variances, common pairwise correlation  $\rho \geq 0$ ; Kish (1965). 60%: [Kim et al., ICML 2025](#) (HELM, MMLU). 2.2: [Kohli, arXiv:2605.29800 \(2026 preprint\)](#), 9 frontier judges, 7 families. Amdahl, AFIPS SJCC 1967.

# More forks are not more witnesses

100 forks that share a little: 9 witnesses. A lot: 2. Identical: 1.

$$\text{Var}(\bar{Y}_N) = \sigma^2 \left( \rho + \frac{1 - \rho}{N} \right)$$


$$N_{\text{eff}} = \frac{N}{1 + (N - 1)\rho}$$

$\rho$ : how much the copies share.

A little = 0.1, a lot = 0.5, identical = 1.

4 re-runs  $\times$  100 tests = 400 executions, 100 evidence IDs. Pool within a candidate; compare across.

Different models, both wrong: same answer 60% (chance 33%). 9 LLM judges: 2.2 independent votes.



**Amdahl:** speedup =  $N / (1 + (N - 1)f)$ ; 100 cores, 10% serial:  $9.2\times$ .  $\rho$  is the serial fraction of your evidence.

🔗 **Forks multiply executions, not evidence.**

Equal variances, common pairwise correlation  $\rho \geq 0$ ; Kish (1965). 60%: Kim et al., ICML 2025 (HELM, MMLU). 2.2: Kohli, arXiv:2605.29800 (2026 preprint), 9 frontier judges, 7 families. Amdahl, AFIPS SJCC 1967.

# The runtime holds some of the cluster labels

## Methods for dependent evidence

|                         |                      |
|-------------------------|----------------------|
| design effect           | Kish 1965            |
| cluster-robust variance | Liang & Zeger 1986   |
| dependent effect sizes  | Hedges et al. 2010   |
| lineage as annotation   | Green et al. PODS'07 |

All of them need cluster labels: who shares what.

## What the runtime can and cannot record

- ✓ shared parent, seed, test, fixture
- ✓ the model each child called: the egress gateway sees it
- ✗ blind spots shared even across different models (Kim et al.)

**ANALOGY** My co-authors and I had to infer this sharing from outside. Address linking tied most bitcoin mined before February 2011 to 64 agents; malignant cells cluster by patient tumour, not by cell type. A fork runtime records who shares what with every result, like a family tree.


**Others must infer the clusters. A fork runtime can record its own.**

Blackburn et al., [arXiv:2206.02871](#) (2022); Saurty-Seerunghen et al., [iScience](#) 2026; Eliaz, [arXiv:2607.09689](#) §5.  
Crossed factors: Cameron, Gelbach & Miller, [JBES](#) 2011 (multiway); [REStat](#) 2008 (wild-cluster bootstrap).

# Fork is also how you run the control

For a comparison, positive covariance helps.

$$\text{Var}(Y_A - Y_B) = \text{Var}(Y_A) + \text{Var}(Y_B) - 2 \text{Cov}(Y_A, Y_B)$$



**PROPOSED** The seed policy belongs in the fork API: `seed=copy|reseed` (slide 30).


Test 2 uses this: siblings and strangers share slot and host, so only ancestry differs (slide 45).

**Share randomness to compare; vary it to corroborate.**

Common random numbers: [Glasserman & Yao, Management Science \(1992\)](#).


# When does a swarm gel?

## Execution ancestry




Which state produced this child?

## Evidence flow



Which observations influenced a decision?

## Composition



Which sets need joint validation?


**ANALOGY · PHYSICS** Stockmayer (J. Chem. Phys. 11:45, 1943): enough cross-links join branched chains into one gel. Enough evidence edges can join separate lineages into one cluster.

**Ancestry is a tree. Evidence flow is not.**

Stockmayer (1943): a precedent under idealised assumptions; an agent threshold must come from the real dependency model.

# Two green patches can fail together

$$\Gamma_{ij} = U(\{i, j\}) - U(\{i\}) - U(\{j\}) + U(\emptyset)$$



Illustrative interaction graph

- $U(A)$ : one fixed objective for change set  $A$ , same parent.
- 8 candidates: 28 pair tests, 256 subsets.
- Effects among 3 or more patches can pass every pair test.
- Rebuild and evaluate every chosen composition (contract condition 7, end matter: verify the merged result).

**ANALOGY · BIOLOGY** Genetics calls a strongly negative  $\Gamma$  synthetic lethality. My co-authors and I screened the genome for genes synthetically lethal with an SRC-3 inhibitor (Commun. Biol. 4:399, 2021).

**Rebuild and evaluate the composition you ship.**

# Every worker returns a receipt, not a score

$$r_k = (\hat{\theta}_k, J_k, n_k, \mathcal{E}_k, \mathcal{L}_k, m_k)$$

result  $(\hat{\theta}, J, n)$ , evidence  $(\mathcal{E})$ , origin  $(\mathcal{L}, m)$

**BUILT**


- numeric summaries merge in any tree order
- count evidence once: a repeated ID stops the merge
- lineage travels with the result; nothing consumes it yet

**PROPOSED**

- one execution per evidence ID: a retry or hedge is never counted twice
- sharing unknown: say “I don’t know” (abstain)

*“Worker count by itself carries no statistical information.”*

arXiv:2607.09689, reduction invariant



duplicate  $e_1$ : merge stops

**MEASURED** 4-worker trace: 6.70 s end to end; pooled 4.942 vs. full-sample 4.945.


## The estimator is Cochran’s. The contract is what was missing.

Eliaz, [arXiv:2607.09689](https://arxiv.org/abs/2607.09689) (2026), §2–5 (trace: §4); code: [github.com/zo123/boltzmann-mapreduce](https://github.com/zo123/boltzmann-mapreduce); Cochran (1954).  
 Replay: DeepSeek-V4 §5.2.5 (DSec).

# A sandbox bounds actions, not claims


## Honest, unequal shards

Five logistic shards,  $n = 60$  to 5,000. Distance to the full-data MLE; lower is better.



## One shard forges its precision

A distant estimate from a 2,000-point shard, precision inflated a further 50×.



Mean  $\pm$  SD over 8 fixed seeds (not CIs); sample-size weighting not yet compared. The heuristic is not a Byzantine guarantee.


**PROPOSED** A copy can overstate its precision ( $n \cdot J$ ): re-run or re-score it on held-out data the controller holds.

## Precision is an input the child should not control.

Synthetic checks, Eliaz, [arXiv:2607.09689](https://arxiv.org/abs/2607.09689), §4. Not agent-task benchmarks.

# Test 1: does restore beat a warm cache?

I am wrong if the whole interval for the gain is under 10s, at any  $N$ .




**A result where the build cache wins is a good result for this room.**

Pre-registration, tag prereg-v1, [github.com/zozo123/cam-talk-london-26](https://github.com/zozo123/cam-talk-london-26). Restore-cost practice: REAP (ASPLOS'21), FaaSnap (EuroSys'22). Decision rules in full: end matter.

# Why I expect siblings to fail together

In simulated actomyosin, the Arp2/3 level changes collective behaviour.



schematic; gold = Arp2/3

| Arp2/3 level | Reported network response                                                                                             |
|--------------|-----------------------------------------------------------------------------------------------------------------------|
| High         | stalls                                                                                                                |
| Low          | contracts                                                                                                             |
| Intermediate | loosely connected clusters may collapse suddenly in motor-driven avalanches, reminiscent of experimental “cytoquakes” |

We forecast these avalanches with network science and machine learning (Li et al., JPCB 2021).


## A motivation, not a prediction.

Analogy: these results do not predict software behaviour.

Team results: [Liman et al. \(PNAS, 2020\)](#); [Li et al. \(JPCB, 2021\)](#); [Eliaz et al. \(PRE, 2020\)](#) (linker valency).

## Test 2: do siblings fail together?

I am wrong if the sibling excess  $\Delta\rho$  is confidently below 0.05.



$\Delta\rho$ : how much more siblings fail together than strangers.

Siblings: technical replicates. Strangers: independent experiments, matched on slot and host.

Why 0.05: that much extra correlation takes 9 siblings down to 6.4 witnesses.

**Each rejection rule was written before the data.**

Pre-registration, tag prereg-v1. Replicates: Gilad, Eliaz et al., Commun. Biol. 2021. Decision rules in full: end matter.

# What two tests cannot separate

A win still needs a mechanism.


| Vary                           | Observe                             | Question                      |
|--------------------------------|-------------------------------------|-------------------------------|
| Prefix cost and dirty pages    | Resource use and memory             | When does reuse pay?          |
| No sharing / archive / routing | Quality at equal full cost          | When does communication help? |
| Evidence overlap and coupling  | Calibration and error               | Does the reducer overcount?   |
| Evaluation budget              | False acceptance and time to target | Where should compute go?      |
| Composition order and size     | Joint success and regressions       | Which interactions matter?    |

A fast fork with a worse search policy can still lose.

## Separate a runtime gain from a search-policy or evaluation advantage.

Ablation plan (proposed). Tries as part of the result: Stroebl, Kapoor & Narayanan, arXiv:2411.17501.

# The same work order, run as it should be



Review 2 would not count as a second witness.

**Search can race. Promotion must not.**

A design, not a run. Why three: when a wrong answer costs more than none, the best number of tries is often under ten (Stroebl et al. 2024). Hidden verifiers, isolated from agents: Kimi K3 report §4.2.6.

# Open problems where this room is ahead of me

1. **Dependence from lineage and evidence edges.** Estimate  $\rho$ , how much forks share, from fork trees and evidence manifests.
2. **Leases that fork.** Capabilities re-minted on restore, revocable, never ambient.
3. **Warm state and side channels.** Pack forks by trust, not only by utilisation.
4. **Evidence commit.** Held effects for remote services and for people.
5. **Layers or VM fork.** When does a content-addressed build beat a live fork?
6. **Human-gate latency.** Snapshot, destroy and restore across a day-long wait.

I would like help with the first two.

# The hedging paper already had the answer

*“a particular request exercises an untested code path”*

Dean & Barroso, The Tail at Scale, CACM 2013,  
“Canary requests”

- ✗ crashes or long delays on thousands of servers at once
- ✓ sent to one or two leaf servers first; the rest only if the canary succeeds

A canary for correlation: a few siblings and strangers before a hundred forks (proposed).  
Agreement is not proof: 9 frontier judges who all agreed were still wrong 9.1% of the time (Kohli).

## Hedge for slow replicas. Canary for shared failures.

Dean & Barroso, CACM 56(2):74–80, 2013, p. 79. Our Airflow Factory runs, Sep 2026: reviews 1 and 2. Kohli, [arXiv:2605.29800 \(2026 preprint\)](#).

**Our first run**  
new code, no test of its own



nine green forks would  
share the gap (hypothetical)

Hedging works when failures are independent.

# Contributions

**BUILT** a receipt contract: reused evidence stops the merge; lineage travels with every result

**MEASURED** two real factory runs

**PROPOSED** H1: restore beats a warm cache by  $>10$  s

**PROPOSED** H2: siblings from one snapshot fail together more than strangers

2009: else: VM\_wait(ID)

2026: else: reduce(receipts)  $\rightarrow N_{\text{eff}}$ , or abstain (proposed)

 **Forks multiply executions, not evidence.**

Help wanted: estimating  $\rho$  from lineage; leases that fork.

arXiv:2607.09689

# Fork lineage, in full

|             | System                                                                   | What it forks or moves                                                                                                                                             | Reported                                                                                                                                                                                          |
|-------------|--------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| 2003        | Xen (Cambridge)                                                          | paravirtualised VMs                                                                                                                                                | base of the Xen rows below                                                                                                                                                                        |
| 2005        | Live migration (Cambridge)                                               | pre-copied memory, to another host                                                                                                                                 | 60 ms downtime, Quake 3 server                                                                                                                                                                    |
| 2005        | Potemkin                                                                 | honeypot VMs, CoW memory                                                                                                                                           | 521 ms clone to first reply                                                                                                                                                                       |
| 2009        | SnowFlock                                                                | VM_fork(N) across hosts, on Xen                                                                                                                                    | 600–800 ms                                                                                                                                                                                        |
| 2020        | Catalyzer                                                                | sfork of a running gVisor sandbox                                                                                                                                  | under 1 ms, best case                                                                                                                                                                             |
| 2023        | Nephele; MITOSIS                                                         | unikernel VMs on Xen; fork over RDMA                                                                                                                               | —                                                                                                                                                                                                 |
| <b>2026</b> | DeltaBox<br>Shepherd<br>Crab<br>Planarian<br>Kimi K3 AgentENV (training) | CoW files + incremental CRIU<br>overlay + backend checkpoints<br>eBPF effect classes per turn<br>local and remote statepoints<br>Firecracker pause, fork, snapshot | 10.83 ms checkpoint<br>fork 134–143 ms, flat from 42 MB to 5.8 GB<br>over 75% of turns need no checkpoint<br>remote rollback by compensating actions<br>checkpoint as low as 133 ms, resume 49 ms |

SOSP'03; NSDI'05; SOSP'05; EuroSys'09; ASPLOS'20; EuroSys'23; OSDI'23; [DeltaBox](#); [Shepherd](#); [Crab](#); [Planarian](#); [Kimi K3](#) §5.3.2.

# Fork cost model and API timings

## Cost model PROPOSED

$$C_{\text{cold}} = N(P + W + V)$$

$$C_{\text{fork}} = P + H + N(R + D + W + V)$$

$$(N - 1)P > H + N(R + D)$$

prepare  $P$ , capture  $H$ , restore  $R$ , diverge  $D$ , work  $W$ , evaluate  $V$ .  $W$  includes model calls and cancels. Each child still rebuilds its patch.

## Three API paths MEASURED

|             | Operation               | Success   | p50    | p95    |
|-------------|-------------------------|-----------|--------|--------|
| Daytona     | create                  | 1024/1024 | 0.20 s | 1.12 s |
| Tensorlake  | create                  | 256/256   | 3.44 s | 9.00 s |
| My platform | restore + run + capture | 255/256   | 6.87 s | 9.04 s |

My platform: 141 MB named snapshot, concurrency 12. Different operations: not a ranking. Teardown excluded. Memory capture unverified, so that row is restore fan-out.

Eliaz, [arXiv:2607.09689](https://arxiv.org/abs/2607.09689), Table 2.

# The fork contract, in full

## Fork contract, written before the fork

1. fork from an immutable, content-addressed parent
2. each branch: own lease, TTL, budget, provenance
3. no publishing credentials in a coding cell
4. record parent checkpoint, branch ID, evidence IDs
5. **never count siblings that reuse the same evidence as independent agreement**
6. merge deterministically; conflicts are evidence
7. full verification and review on the merged result
8. tear down every child, on any exit    our second run

Airflow Factory fork contract (design doc). Slide 31 shows conditions 5 and 8.

# Pre-registered decision rules

## Test 1: does restore beat a warm cache?

| Step                | Design                                                                                                                                                                                                                                    | Decision                                                                                                                                                                        |
|---------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Gate 0              | a RAM nonce, snapshot, restore 3 children: does it survive?                                                                                                                                                                               | if not, every result is “restore fan-out”                                                                                                                                       |
| Fidelity<br>H1 cost | deterministic tests, restored vs. direct<br>cached template vs. restore, cold as reference;<br>$N \in \{3, 6, 12\}$ ; 20 interleaved reps per arm and $N$ ;<br>makespan to first test results, one clock for both<br>arms; no model calls | any output mismatch voids the arm<br>supports if the 95% CI of the median gain is<br>above 10 s at every $N$ ; rejects if the 98.3%<br>(Bonferroni) CI is below 10 s at any $N$ |

## Test 2: do siblings fail together?

| Step        | Design                                                                                                                                                                                           | Decision                                                                                                                                                       |
|-------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------|
| H2 coupling | 12 snapshot families $\times$ 3 siblings, each paired with 3<br>strangers restored from other families, same slot,<br>same host; the race test in lockstep rounds, at least<br>300 per condition | supports if a one-sided sign-flip test gives<br>$p < 0.05$ and $\Delta\rho > 0.05$ ; rejects if the 90% CI<br>of $\Delta\rho$ is below 0.05; else inconclusive |
| The loop    | repair 1 re-sampled 9 times from its recorded input;<br>the only model calls                                                                                                                     | descriptive; predicted: $\geq 5$ of 9 repeat the<br>out-of-plan edit                                                                                           |

Pre-registration in [github.com/zozo123/cam-talk-london-26](https://github.com/zozo123/cam-talk-london-26), tags prereg-v1 to prereg-v3.

# Research record: biophysics

PhD, University of Houston, with the Rice Center for Theoretical Biological Physics

**Eliaz**, Nedelec, Morrison, Levine & Cheung. Insights from graph theory on the morphologies of actomyosin networks with multilinkers. *Phys. Rev. E* 102:062420 (2020). First author.

Liman, . . . , **Eliaz**, . . . , Cheung. The role of the Arp2/3 complex in shaping the dynamics and structures of branched actomyosin networks. *PNAS* 117:10825 (2020).

Li, Liman, **Eliaz** & Cheung. Forecasting avalanches in branched actomyosin networks with network science and machine learning. *J. Phys. Chem. B* 125:11591 (2021).

Zegarra, . . . , **Eliaz**, . . . . Impact of hydrodynamic interactions on protein folding rates depends on temperature. *Phys. Rev. E* 97:032402 (2018).

Zhang, Nde, **Eliaz**, Jennings, Cieplak & Cheung. CaXML: Chemistry-informed machine learning explains mutual changes between protein conformations and calcium ions in calcium-binding proteins using structural and topological features. *Protein Science* 34:e70023 (2025).

## Thesis

**Eliaz**. Physics of self-assembly in complex matter. PhD thesis, University of Houston (2020); self-published in book form as *Physics of Complex (Biological) Matter* (2023).

Full author lists and positions: CLAIM-FENCE.md. Meeting abstracts (Biophysical Society, APS March Meeting 2018–2021) are not listed as papers.

# Research record: genomics

## Genome architecture and bioinformatics

Harris, . . . , **Eliaz**, . . . , Lieberman Aiden, Rowley. Chromatin alternates between A and B compartments at kilobase scale for subgenic organization. *Nat. Commun.* 14:3303 (2023).

Kaushal, . . . , **Eliaz**, . . . . CTCF loss has limited effects on global genome architecture in *Drosophila* despite critical regulatory functions. *Nat. Commun.* 12:1011 (2021).

Barkai, . . . , **Eliaz**, . . . . OffRisk: a docker image for annotating CRISPR off-target sites in the human genome. *Bioinformatics Advances* 3:vbad138 (2023).

## Cancer genomics

Saurty-Seerunghen, El-Habr, **Eliaz**, . . . , Junier. A unique malignant cell type per patient tumor encoded in each cancer cell transcriptome. *iScience* 29:115139 (2026).

Gilad, **Eliaz**, . . . . A genome-scale CRISPR Cas9 dropout screen identifies synthetically lethal targets in SRC-3 inhibited cancer cells. *Commun. Biol.* 4:399 (2021).

Gilad, **Eliaz**, . . . . Reply to: Target expression is a relevant factor in synthetic lethal screens. *Commun. Biol.* 5:836 (2022).

Gilad, **Eliaz**, . . . . Drug-induced PD-L1 expression and cell stress response in breast cancer cells can be balanced by drug combination. *Sci. Rep.* 9:15099 (2019).

Full author lists and positions: CLAIM-FENCE.md.

# Research record: preprints, systems and talks

## Preprints

**Eliaz**. Evidence-Aware MapReduce for Forkable Compute. arXiv:2607.09689 (2026). Sole author; code [github.com/zozo123/boltzmann-mapreduce](https://github.com/zozo123/boltzmann-mapreduce). Blackburn, Huber, **Eliaz**, . . . , Lieberman Aiden. Cooperation among an anonymous group protected Bitcoin during failures of decentralization. arXiv:2206.02871 (2022). **Eliaz**, Danovich & Gasic. Poolkeh finds the optimal pooling strategy for a population-wide COVID-19 testing. medRxiv (2020). First author. Hitz, . . . , **Eliaz**, . . . , Cherry. The ENCODE uniform analysis pipelines. bioRxiv (2023). The ENCODE Project Consortium (member). The Encyclopedia of DNA Elements. bioRxiv (2026). Chamo & **Eliaz**. Continuation geometry resolves apparent branch splitting in unequal-mass three-body orbits. ai.viXra:2608.0069 (2026); not peer reviewed.

## Systems and open source

Airflow Factory, an open-source Apache Airflow software factory: [github.com/zozo123/ariflow-swfactory](https://github.com/zozo123/ariflow-swfactory). This talk's slides and notes: [github.com/zozo123/cam-talk-london-26](https://github.com/zozo123/cam-talk-london-26). apache/airflow: sandbox toolset with Docker sbx and OpenSandbox backends (merged, 2026); two more backends (open). openclaw/crabbox, nushell, and sandbox PRs to Spark, Kafka Connect and mecatl (2026). The Sandbox Shift: a field manual for running untrusted code (2026).

## Talks, October 2026

Cambridge SRG (this talk), 15 Oct.  
RustChinaConf, Shenzhen: Disposable Runners, Warm Cargo Factory, 17 Oct.  
Zenity AI Agent Security Summit, New York: Your Agent Escaped Without Escaping the Sandbox, 21 Oct.

# References: systems

### Fork, snapshot, isolation

Barham et al. Xen and the art of virtualization. SOSP 2003.  
Vrable et al. Scalability, fidelity, and containment in the Potemkin virtual honeyfarm. SOSP 2005.  
Lagar-Cavilla et al. SnowFlock. EuroSys 2009.  
Madhavapeddy et al. Jitsu. NSDI 2015.  
Manco et al. My VM is lighter (and safer) than your container. SOSP 2017.  
Baumann et al. A fork() in the road. HotOS 2019.  
Agache et al. Firecracker. NSDI 2020.  
Du et al. Catalyzer. ASPLOS 2020.  
Brooker et al. Restoring uniqueness in microVM snapshots. arXiv:2102.12892, 2021.  
Ustiugov et al. Benchmarking, analysis, and optimization of serverless function snapshots (REAP). ASPLOS 2021.  
Lupu et al. Nephele: cloning unikernel-based VMs. EuroSys 2023.  
Clark et al. Live migration of virtual machines. NSDI 2005.

### Leases and fencing

Gray & Cheriton. Leases. SOSP 1989.  
Burrows. The Chubby lock service. OSDI 2006.  
Kleppmann. How to do distributed locking. 2016.

### Agent sandboxes, 2026

Dong et al. DeltaBox. arXiv:2605.22781. Yu et al. Shepherd. arXiv:2605.10913. Wu et al. Crab. arXiv:2604.28138.  
Guo et al. Planarian: managing agent state with statepoints. arXiv:2609.35366. Zheng et al. arXiv:2608.22928.  
Kimi Team (Moonshot AI). Kimi K3: Open Frontier Intelligence. Tech. report, 2026, §4.1.2, §4.2.4, §4.2.6, §5.3.2. DeepSeek-AI. DeepSeek-V4 (DSeC). arXiv:2606.19348.  
Wei et al. No provisioned concurrency: fast RDMA-codesigned remote fork (MITOSIS). OSDI 2023. Ao et al. FaaSnap. EuroSys 2022.

### Effects, authority, provenance

Nightingale, Chen & Flinn. Speculative execution in a distributed file system. SOSP 2005.  
Nightingale et al. Rethink the sync. OSDI 2006.  
Cully et al. Remus. NSDI 2008.  
Watson et al. Capsicum. USENIX Security 2010.  
Watson et al. CHERI. IEEE S&P 2015.  
Birgisson et al. Macaroons. NDSS 2014.  
Muniswamy-Reddy et al. Provenance-aware storage systems. USENIX ATC 2006.  
Pasquier et al. Practical whole-system provenance capture. SoCC 2017.  
Torres-Arias et al. in-toto. USENIX Security 2019.  
O'Callahan et al. Engineering record and replay for deployability. USENIX ATC 2017.

# References: evidence and evaluation

### Dependent evidence

Dorfman. The detection of defective members of large populations. *Ann. Math. Stat.* 14(4):436–440, 1943.

Cochran. The combination of estimates from different experiments. *Biometrics* 1954.

Kish. *Survey Sampling*. Wiley, 1965.

Amdahl. Validity of the single processor approach to achieving large scale computing capabilities. *AFIPS SJCC* 1967.

Liang & Zeger. Longitudinal data analysis using generalized linear models. *Biometrika* 1986.

Hedges, Tipton & Johnson. Robust variance estimation in meta-regression with dependent effect size estimates. *Research Synthesis Methods* 2010.

Green, Karvounarakis & Tannen. Provenance semirings. *PODS* 2007.

Cameron, Gelbach & Miller. Bootstrap-based improvements for inference with clustered errors. *REStat* 2008.

Cameron, Gelbach & Miller. Robust inference with multiway clustering. *JBES* 2011.

Glasserman & Yao. Some guidelines and guarantees for common random numbers. *Management Science* 1992.

Stockmayer. Theory of molecular size distribution and gel formation in branched-chain polymers. *J. Chem. Phys.* 11:45, 1943.

### Shared errors across models

Kim, Garg, Peng & Garg. Correlated errors in large language models. *ICML 2025*, PMLR 267. [arXiv:2506.07962](#).

Goel et al. Great models think alike and this undermines AI oversight. *ICML 2025*. [arXiv:2502.04313](#).

### Evaluation under search

Dean & Barroso. The tail at scale. *CACM* 2013 (hedged requests; also fan-out, good enough, canary requests, mutations).

Dean. Designs, lessons and advice from building large distributed systems (“Numbers everyone should know”). *LADIS keynote*, 2009.

Brown et al. Large Language Monkeys: scaling inference compute with repeated sampling. [arXiv:2407.21787](#), 2024.

Von Arx, Chan & Barnes (METR). Recent frontier models are reward hacking. 5 June 2025.

Zhong, Raghunathan & Carlini. ImpossibleBench. [arXiv:2510.20270](#), 2025.

Dwork et al. The reusable holdout. *Science* 2015.

Stroebl, Kapoor & Narayanan. The limits of inference scaling through resampling (v1: Inference scaling fLaws). [arXiv:2411.17501](#), 2024.

Russo & Zou. Controlling bias in adaptive data analysis using information theory. *AISTATS* 2016.

**Shared errors across models: 2026 preprints**

Jo, Garg & Raghavan. The subjectivity of monoculture. [arXiv:2602.24086](#), preprint.

Kohli. Nine judges, two effective votes: correlated errors undermine LLM evaluation panels. [arXiv:2605.29800](#), preprint.

Chen. When does combining language models help? A co-failure ceiling on routing, voting, and mixture-of-agents across 67 frontier models. [arXiv:2606.27288](#), preprint.

Bone, Stephany & del Rio-Chanona. Monocultural biases: correlated biases in large language models lead to unequal systemic exclusion rates in hiring. [arXiv:2609.22169](#), preprint.

# References: own work and documentation

### Speaker and co-authors

Y. Eliaz. Evidence-Aware MapReduce for Forkable Compute. arXiv:2607.09689, 2026.

[github.com/zozo123/ariflow-swfactory](https://github.com/zozo123/ariflow-swfactory), runs SELFHOST-2/3, 27 Sep 2026.

Blackburn, Huber, Eliaz et al. Cooperation among an anonymous group protected Bitcoin during failures of decentralization. arXiv:2206.02871, 2022.

Saurty-Seerunghen, El-Habr, Eliaz et al. A unique malignant cell type per patient tumor encoded in each cancer cell transcriptome. iScience 29:115139, 2026.

Hitz, . . . , Eliaz, . . . , Cherry. The ENCODE uniform analysis pipelines. bioRxiv 2023, doi:10.1101/2023.04.04.535623.

Gilad, Eliaz et al. A genome-scale CRISPR Cas9 dropout screen identifies synthetically lethal targets in SRC-3 inhibited cancer cells. Commun. Biol. 4:399, 2021.

Barkai, . . . , Eliaz, . . . . OffRisk. Bioinformatics Advances 3:vbad138, 2023.

Eliaz, Danovich & Gasic. Poolkeh. medRxiv 2020.

### Speaker's physics work (slides 7 and 44)

Eliaz et al. PRE 102:062420, 2020; Liman et al. PNAS 117:10825, 2020; Li et al. JPCB 125:11591, 2021.

### Documentation

OBuilder ([github.com/ocurrent/obuilder](https://github.com/ocurrent/obuilder)); day10 ([tunbury.org](https://tunbury.org), 16 Feb 2026); Docker Sandboxes security; Firecracker snapshotting, snapshot-support and prod-host-setup docs; CRIU TCP\_REPAIR; gVisor checkpoint/restore.

### RL and exploration

Ecoffet et al. Go-Explore. Nature 2021.